Maximising Revenue per Driver Hour (RDH) through Demand Optimisation

Introduction:

In the hyper-competitive market of New York City, Uber faces significant operational challenges stemming from its five-borough complexity, notorious traffic, and fluctuating demand. Inefficient driver distribution, prolonged idle times, and inconsistent trip profitability threaten both driver earnings and overall platform health. This analysis will leverage ride data to optimize a key metric: Revenue per Driver Hour (RDH). By identifying demand patterns, balancing driver supply, and minimizing downtime, this study will produce actionable strategies to increase driver income, enhance fleet efficiency, and strengthen Uber's performance in one of the world's most demanding transportation landscapes.

Objective:

The primary objective of this analysis is to leverage Uber's ride data to provide a data-driven framework for optimizing operational efficiency and profitability in NYC. This will be achieved through four key pillars:

- **Maximize Driver Earnings through RDH Optimization:** Analyze the core factors (time, location, trip type) that influence Revenue per Driver Hour (RDH). The goal is to develop strategies that minimize driver idle time and maximize earnings per trip, directly leading to higher driver satisfaction, retention, and loyalty on the platform.
- **Enhance Fleet Efficiency via Intelligent Demand-Supply Balancing:** By identifying recurring demand hotspots, peak hours, and travel patterns, we can enable more strategic driver deployment. This reduces the mismatch between rider demand and driver supply, which not only boosts fleet efficiency but also enhances the customer experience by lowering wait times and ETAs.
- **Develop Dynamic Pricing and Incentive Models:** Understand how RDH varies across the city to build and refine data-driven business strategies. These insights will inform the creation of targeted driver incentives, smarter surge pricing algorithms, and fare adjustments that improve platform revenue, especially during periods of low demand or high supply.

- **Empower Strategic, Data-Backed Decision-Making:** Synthesize all analytical findings to provide a clear, actionable foundation for Uber's leadership. These insights will support high-level decisions regarding fleet management, driver incentive programs, fare structures, and potential market expansion, ensuring a sustainable competitive advantage.

Business Impact:

This project delivers tangible business value by transforming raw data into strategic action. The primary impacts are:

- **Boosts NYC Revenue:** Optimizes five-borough fleet deployment and pricing to increase profitability and driver RDH.
- **Reduces Driver Churn:** Delivers higher, more consistent pay to improve driver retention in a competitive market.
- **Builds Rider Preference:** Wins loyalty from taxi and transit users by providing more reliable service and faster pickups.
- **Secures Market Leadership:** Uses data to outperform rivals like Lyft and taxis with superior operational efficiency.

Dataset Overview:

Dataset Name: Uber Ride Data – NYC

Number of Rows: 10,000,000

Number of Columns: 17

Description: This dataset offers a comprehensive, multi-dimensional view of each trip, capturing critical variables such as temporal (timestamps, duration), spatial (geolocations), and financial (itemized fares, payment types) data. With each row representing a distinct journey, this granular structure is ideal for dissecting demand patterns, modelling fleet efficiency, and pinpointing the key drivers of Revenue per Driver Hour (RDH).

- **vendor_id:** A code indicating the TPEP provider that provided the record
- **pickup_datetime:** The date and time when the meter was engaged.
- **dropoff_datetime:** The date and time when the meter was disengaged.
- **passenger_count:** The number of passengers in the vehicle (driver-entered).
- **trip_distance:** The elapsed trip distance in miles reported by the taximeter.
- **rate_code:** The final rate code in effect at the end of the trip
- **store_and_fwd_flag:** Flag indicating if the trip record was stored in vehicle memory before being sent to the vendor.
- **payment_type:** Numeric code signifying how the passenger paid for the trip.

- **fare_amount:** The time-and-distance fare calculated by the meter.
- **extra:** Miscellaneous extras and surcharges (e.g., rush hour or overnight charges).
- **mta_tax:** $0.50 MTA tax automatically triggered based on the metered rate.
- **tip_amount:** Tip amount (auto-recorded for card payments, excludes cash tips).
- **tolls_amount:** Total amount of all tolls paid during the trip.
- **imp_surcharge:** $0.30 improvement surcharge assessed at flag drop (started in 2015).
- **total_amount:** The total amount charged to passengers (excluding cash tips).
- **pickup_location_id:** TLC Taxi Zone where the trip started (meter engaged). **dropoff_location_id:** TLC Taxi Zone where the trip ended (meter disengaged).

KPI Tree:

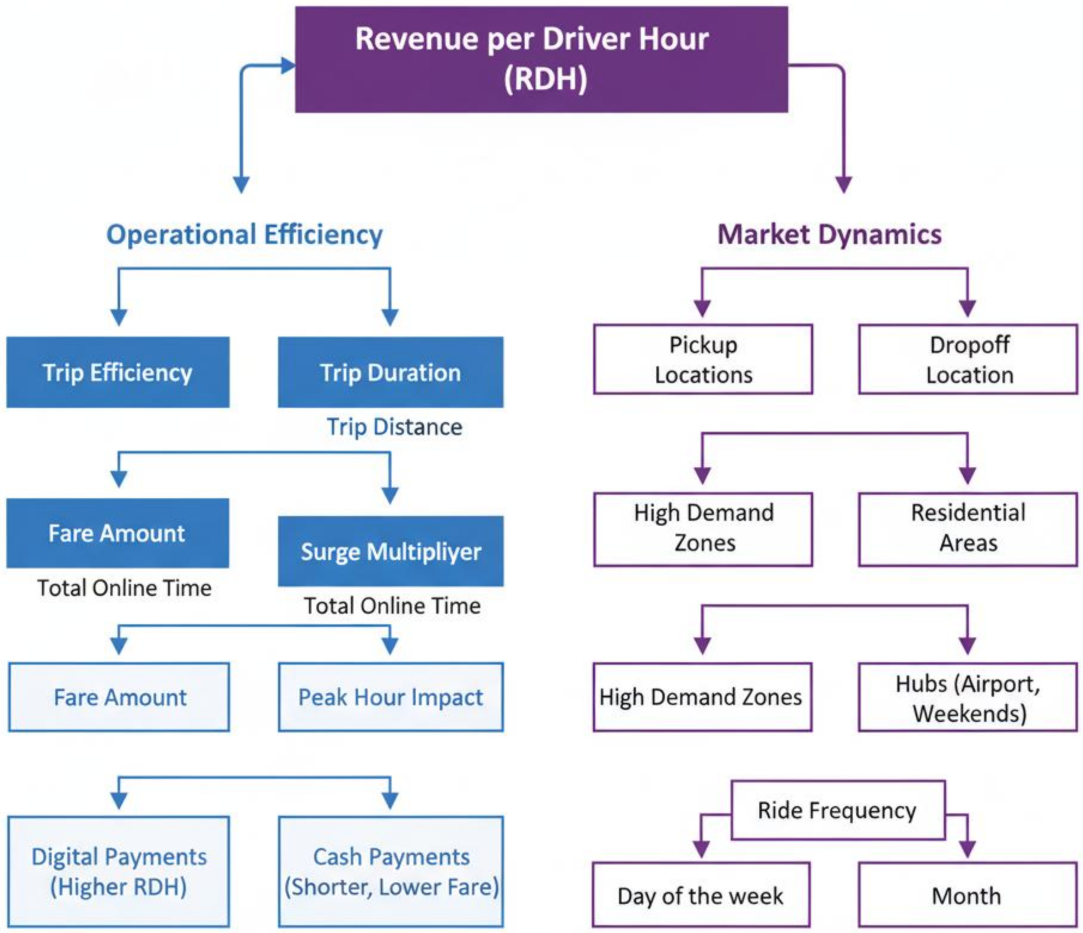

Data Cleaning and Preparation:

To ensure the accuracy, consistency, and reliability of the findings, the raw dataset underwent a comprehensive cleaning and preparation process. Each step was critical in transforming the data into a structured and analysis-ready format suitable for a detailed examination of ride demand, revenue, and driver efficiency.

1. **Data Import and Environment Setup:** The analytical process commenced with the importation of the raw Uber dataset (.csv format) into a Google Colab environment. This platform was selected for its robust capabilities in handling large-scale data processing and analysis.
2. **Handling of Null Values:** An initial inspection revealed the presence of null values in critical columns such as passenger_count and payment_type. To maintain the integrity of subsequent calculations and avoid potential errors in fare analysis or RDH computation, all rows containing these missing values were systematically removed from the dataset.
3. **Removal of Duplicate Records:** The dataset was audited for duplicate entries to ensure each record was unique. From the initial 10,000,000 rows, a total of 607,571 redundant entries were identified and purged. This step refined the dataset to 9,392,429 unique trip records, preventing inflated counts and skewed analytical results.
4. **Feature Engineering: Creation of New Analytical Columns:** To facilitate deeper analysis, new features were engineered from the existing data:

   - **trip_duration**: A crucial metric for efficiency analysis, this column was created by calculating the difference between dropoff_datetime and pickup_datetime, with the result expressed in minutes.
   - **Revenue per Driver Hour (RDH)**: The primary performance indicator for this project was computed by dividing the total_amount of a trip by its trip_duration (converted to hours). This metric is essential for evaluating driver profitability relative to time invested.

5. **Data Enrichment: Modifying Existing Columns:** To enhance the clarity and interpretability of the analysis, columns containing numerical codes—such as pickup_location_id, dropoff_location_id, and payment_type—were transformed. These numeric IDs were mapped to their corresponding meaningful text labels (e.g., specific location names, payment methods like "Credit Card"), making the dataset more intuitive for analysis and visualization.
6. **Outlier Detection and Removal:** An outlier detection process was implemented to prevent extreme or erroneous values from skewing the results. The Interquartile Range (IQR) method, using the 10th and 90th percentiles as thresholds, was applied to key numerical columns like fare_amount, trip_distance, and total_amount. This resulted in the removal of approximately

13.66% of rows, ensuring the analysis is based on a representative and stable subset of the data.

7. **Final Data Consistency Checks:** As a final step, all column names, data formats, and categorical values were standardized. This harmonization ensures smooth data aggregation and accurate KPI calculations (e.g., trip counts, average RDH) across different locations and time periods.

Upon completion of these steps, the dataset was confirmed to be clean, structured, and robust, providing a reliable foundation for the subsequent exploratory data analysis and insight generation.

Exploratory Data Analysis (EDA):

Analysis of Temporal Demand Patterns

Hypothesis: Trip demand in NYC follows predictable daily and weekly patterns, characterized by distinct commuter-driven peaks and differing volumes between weekdays and weekends.

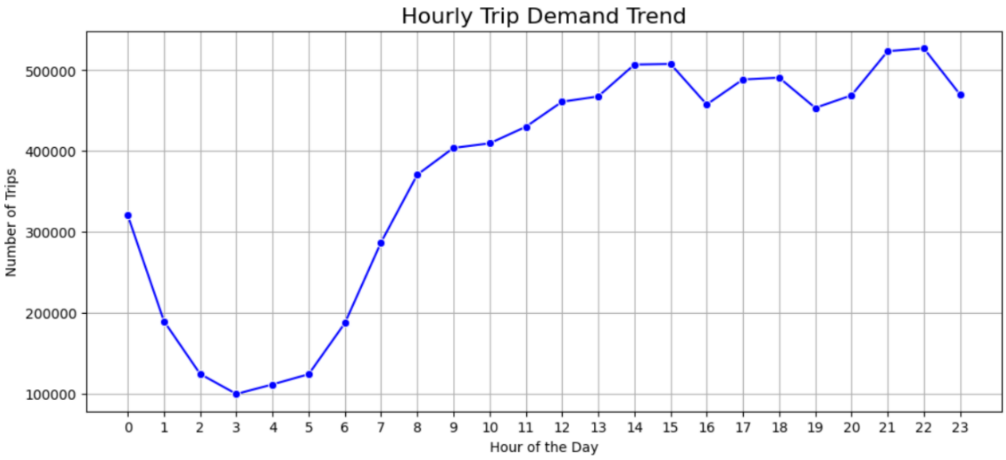

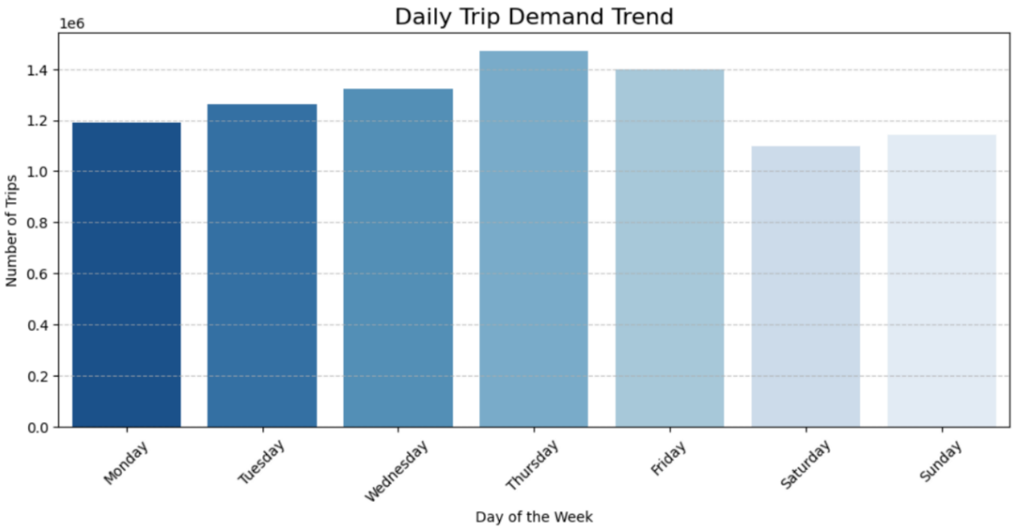

Observations:

- Trip demand follows a distinct dual-peak pattern, surging during the morning commute (7-10 AM) and reaching its highest volume in a sustained evening peak (5-10 PM).
- Ride volume is heavily concentrated on weekdays, peaking on Thursday and Friday, before dropping significantly on the weekend, indicating commuter traffic is dominant.

Insights & Recommendations:

- Launch pre-peak driver bonuses (around 6 AM and 4 PM) to increase driver supply just before the largest demand surges.
- Guide drivers via in-app messaging that late-week evenings (Thurs/Fri, 5-10 PM) offer the highest volume of trips and earning opportunities.

Analysis of Hourly Profitability Trend:

Hypothesis: Revenue per Driver Hour (RDH), a key measure of driver profitability, does not directly mirror trip demand. It is highest during hours where demand significantly outstrips active driver supply, leading to higher fares and surge pricing.

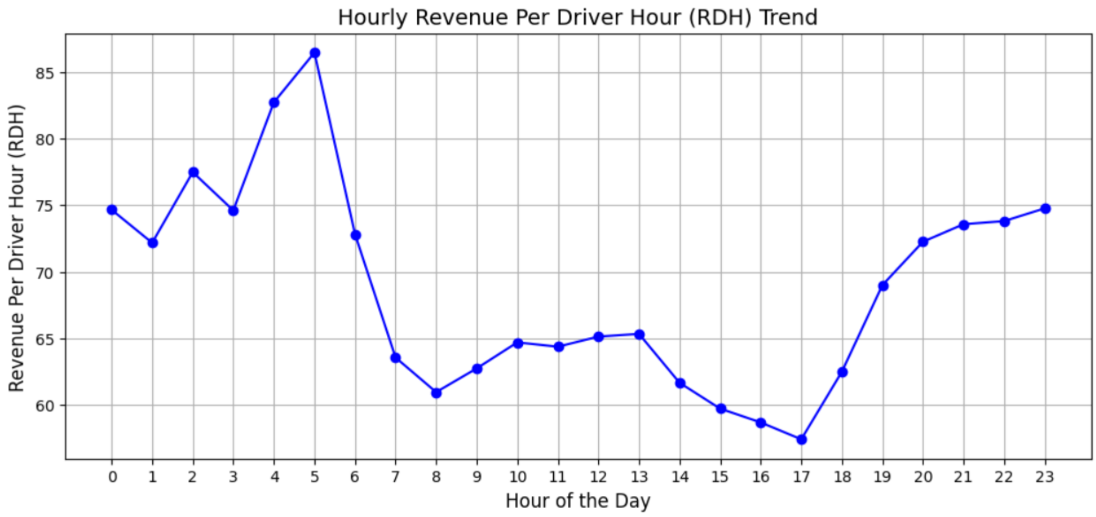

Observations:

- Driver profitability (RDH) peaks in the early morning (3-5 AM) when ride volume is lowest, driven by low driver supply and high-value trips.
- The high-demand morning commute (8-9 AM) corresponds to the lowest RDH of the day, likely due to market saturation and heavy traffic.

Insights & Recommendations:

- Promote a "split-shift" strategy, guiding drivers to work the profitable early morning and late evening hours to maximize their RDH.

- Launch "Night Owl" incentives to ensure driver availability for the high-value, low-supply window between 11 PM and 6 AM.

Analysis of Trip Distance Impact on Fare and Profitability:

Hypothesis: While longer trip distances result in higher total fares, they do not necessarily lead to higher driver profitability (RDH). There is likely an optimal range of trip distances that maximizes a driver's hourly earnings.

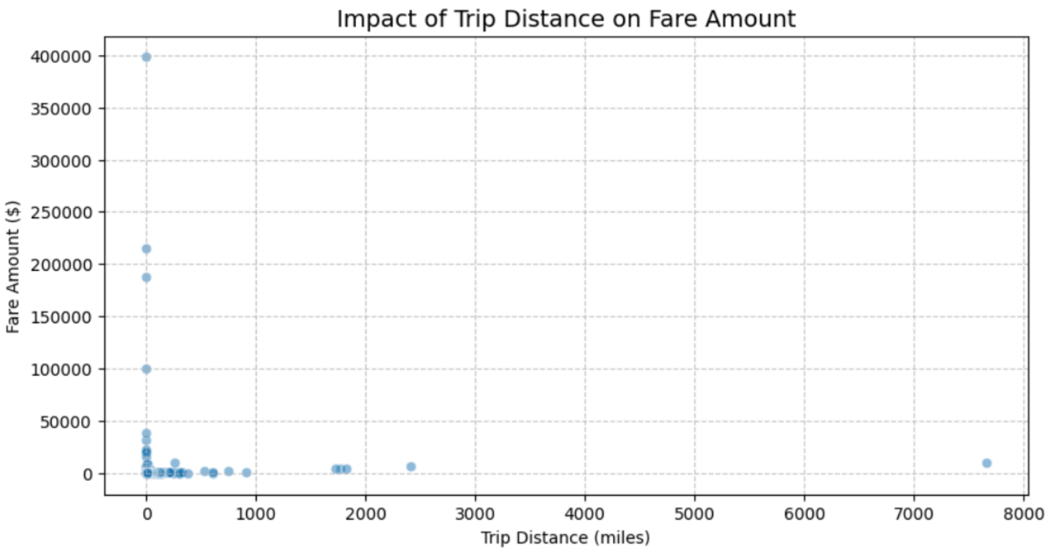

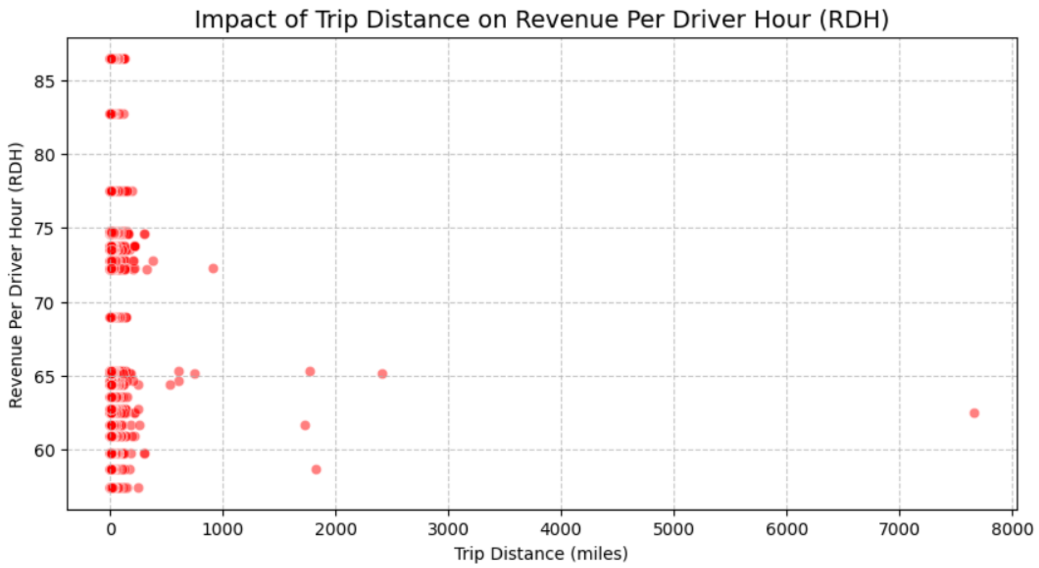

Observations:

- The most profitable trips, measured by RDH, are concentrated in the short-to-medium distance range.
- Very long-distance trips, despite high total fares, yield low hourly earnings (RDH), making them inefficient for drivers.

Insights & Recommendations:

- Adjust the matchmaking algorithm to prioritize a queue of high-RDH short-to-medium distance trips for drivers.
- Re-evaluate the fare structure for long-haul trips to better compensate drivers for unpaid return ("deadhead") miles.

Analysis of Revenue by Payment Type:

Hypothesis: A single payment method is overwhelmingly dominant in generating revenue, indicating a strong preference among NYC riders for one type of transaction over others.

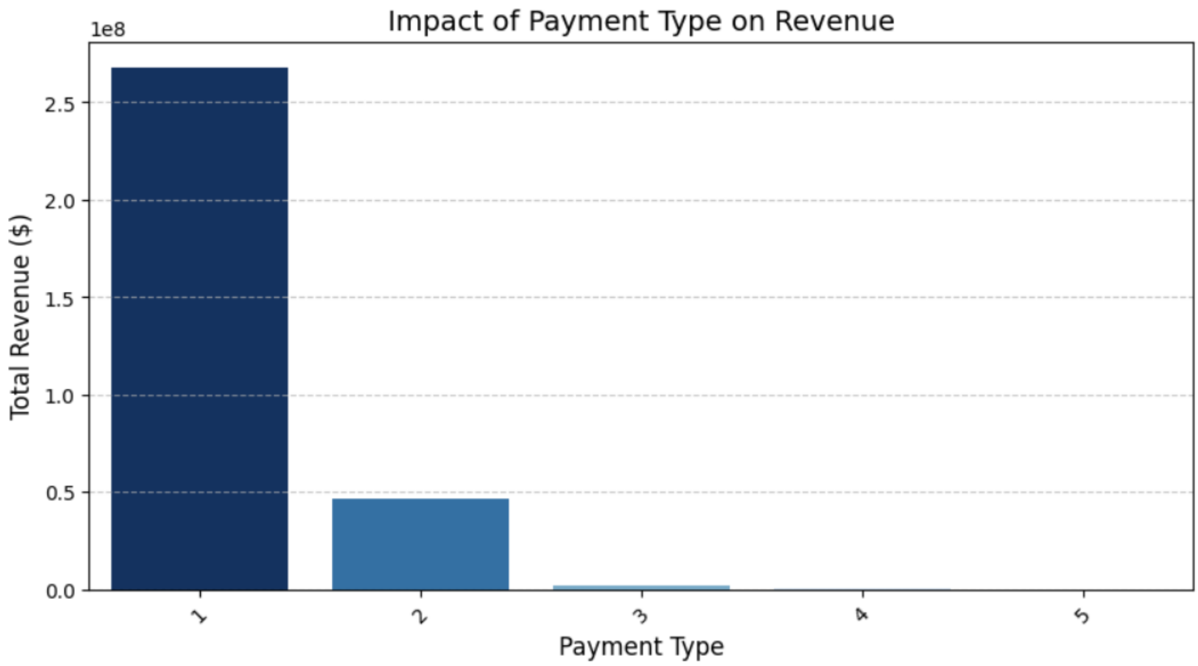

Observations:

- Digital payments (assumed Type 1) are the overwhelmingly dominant source of revenue for the platform.
- Cash payments (assumed Type 2) represent the second-largest, yet significantly smaller, revenue stream.

Insights & Recommendations:

- Double-down on the preferred user behavior by enhancing and expanding in-app digital payment integrations (e.g., UPI, wallets).
- Investigate "No Charge" and "Disputed" trips to identify and resolve operational issues or fraud that cause revenue leakage.

Analysis of Demand vs. Profitability:

Hypothesis: Driver profitability (RDH) often has an inverse relationship with peak trip demand, especially during congested commuter hours where high driver supply and traffic negatively impact earnings efficiency.

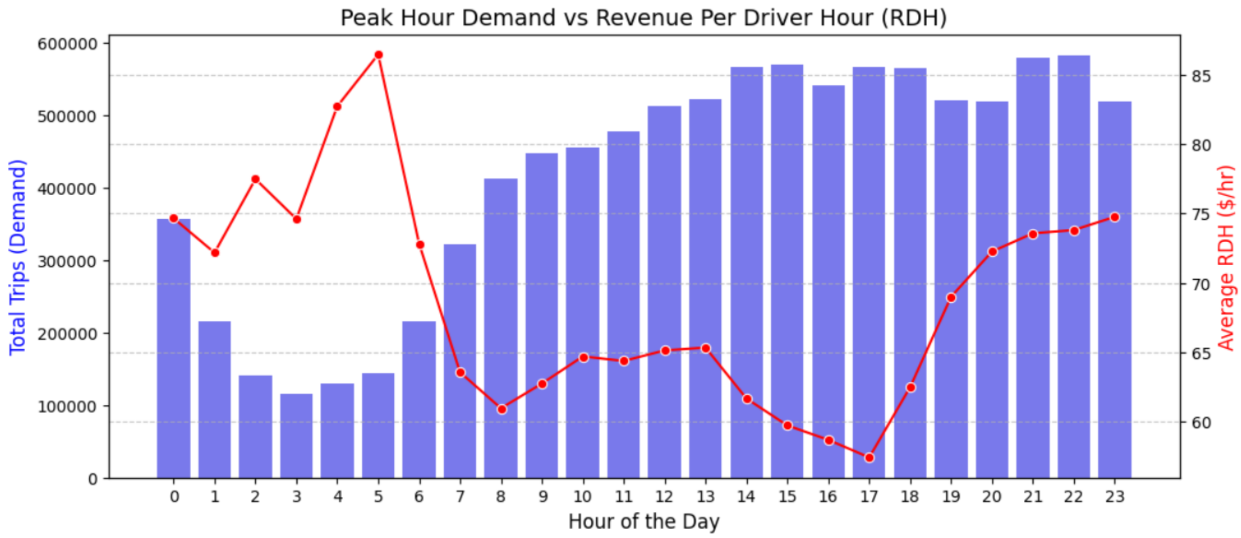

Observations:

- Total ride revenue is highly concentrated in a handful of key pickup and drop-off locations.
- The top pickup zones are different from the top drop-off zones, revealing major asymmetric commuter corridors.

Insights & Recommendations:

- The first priority is to map location IDs to their actual names to make these spatial insights actionable for strategic planning.
- Analyze the specific travel corridors between top zones to predict demand and strategically position drivers ahead of peak commuter flows.

Analysis of Top Earning Locations:

Hypothesis: A select few key locations—likely major commercial hubs, tech parks, or transit centres—are the primary drivers of total revenue, exhibiting distinct patterns for pickups and drop-offs.

Observations:

- A clear inverse relationship exists during the morning commute, where a surge in trip demand causes a sharp drop in driver RDH.
- The most profitable hours for drivers are consistently the ones with the lowest overall ride volume, highlighting the impact of supply and demand imbalance.

Insights & Recommendations:

- Focus driver incentives on the "profitability shoulders" (3-6 AM and post-9 PM) to secure supply during high-value hours.
- Use in-app analytics to communicate to drivers that high trip volume does not always equal high profitability due to traffic.

Analysis of Location-Specific Profitability Trends:

Hypothesis: While the top pick-up zones generally follow the city-wide hourly RDH pattern, the magnitude of the peaks and troughs varies significantly by location, indicating that local factors like traffic and trip type heavily influence driver profitability.

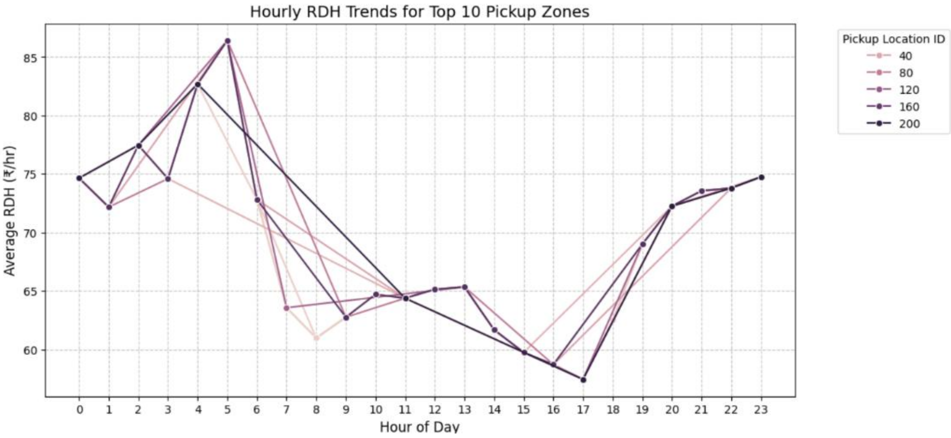

Observations:

- While all top zones follow the same general hourly RDH pattern, the level of profitability varies significantly between them.
- Certain zones (e.g., 160, 200) are clearly more lucrative during the early morning peak, suggesting a higher concentration of valuable trips.

Insights & Recommendations:

- Implement hyper-targeted "Power Zone" bonuses for specific locations during their most profitable hours to maximize ROI.
- Provide drivers with granular, location-aware guidance, advising which zones offer the best earning potential at different times of the day.

Hypothesis Testing

1. Hypothesis: Analysis of RDH by Trip Distance Category

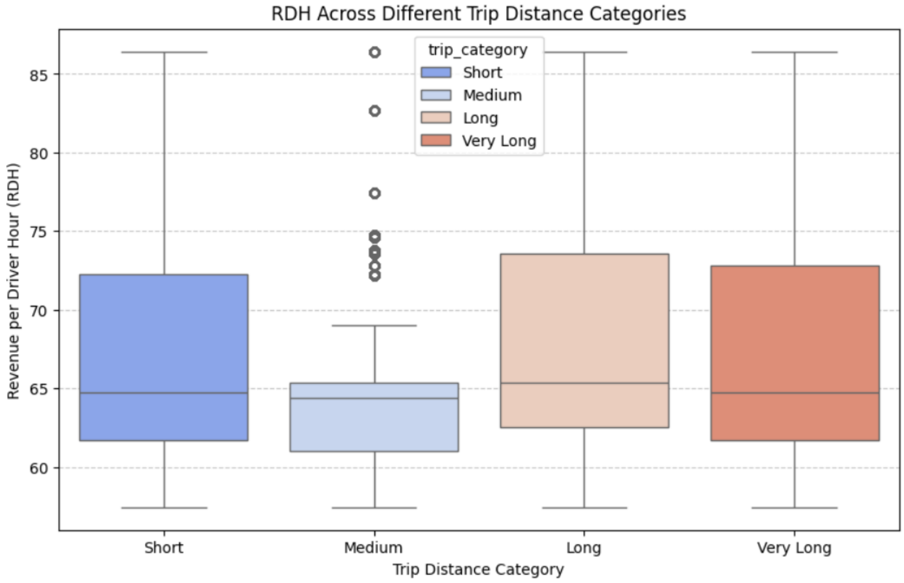

- **Finding: Non-Linear Profitability** Driver profitability (RDH) is not linear with distance. Short and Very Long trips are the most profitable categories, while Medium-distance trips show a distinct dip in hourly earnings.
- **Recommendation: Fix Medium-Distance Fares** Investigate and adjust the fare structure for Medium-distance trips, as this category is the least profitable and likely a significant pain point for drivers.

2. Hypothesis: Analysis of RDH by Peak vs. Non-Peak Hours

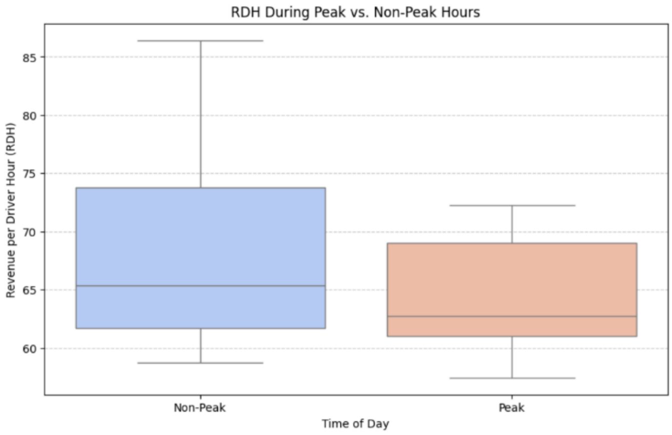

- **Finding: Higher and More Consistent Non-Peak Profitability** Non-Peak hours demonstrate a significantly higher median Revenue per Driver Hour (RDH) and a wider range of high-end earnings. In contrast, Peak-hour earnings are consistently lower and compressed into a tighter range, indicating that factors like traffic and market saturation reliably reduce driver profitability during the busiest times.
- **Recommendation: Incentivize Non-Peak Driving** Launch targeted bonuses for drivers active during historically high-RDH, non-peak periods (especially early mornings and late nights). This ensures a reliable supply for these highly profitable trips and rewards drivers for their availability.

3. Hypothesis: Analysis of RDH by Fare Category

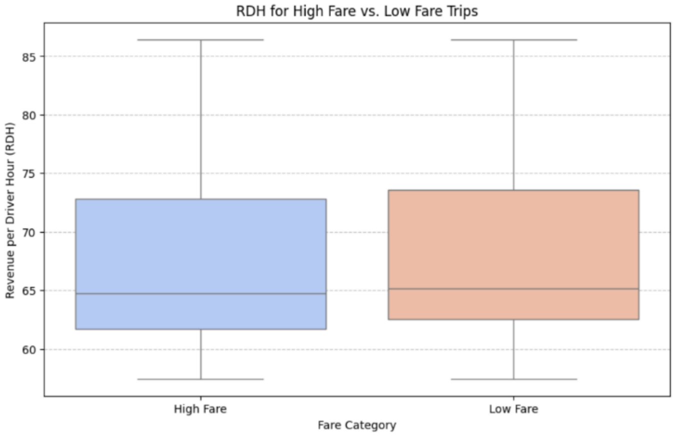

- **Finding: Fare Amount Is Not a Key Driver of RDH** There is no significant difference in the median Revenue per Driver Hour (RDH) between High Fare and Low Fare trips. This indicates that a high absolute fare does not guarantee better hourly earning efficiency, likely because high-fare trips take a proportionally longer amount of time to complete.
- **Recommendation: Prioritize Trip Efficiency over Fare Amount** Shift the focus of the matchmaking algorithm from prioritizing the next "highest fare" trip to prioritizing trips with the highest *predicted RDH*. This means favoring a quick succession of efficient trips over a single, time-consuming high-fare journey.

4. Hypothesis: Analysis of RDH by Day Type

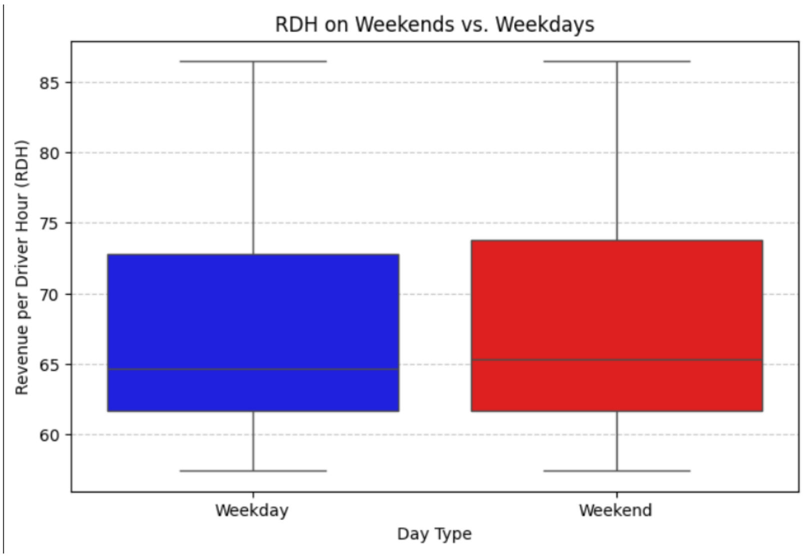

- **Finding: Consistent Profitability Despite Lower Demand** Despite weekends having significantly lower trip volume, the median Revenue per Driver Hour (RDH) is nearly identical to that of weekdays. This suggests that factors like reduced traffic congestion and lower driver supply on weekends effectively compensate for the drop in demand, keeping hourly earning efficiency stable.
- **Recommendation: Promote Weekends as High-Efficiency Driving** Market weekends to drivers as an opportunity for "stress-free earnings." Highlight through in-app communications that they can achieve similar hourly rates as weekdays but with less traffic, which can improve driver satisfaction and retention.

5. Hypothesis: Analysis of RDH by Trip Distance Category

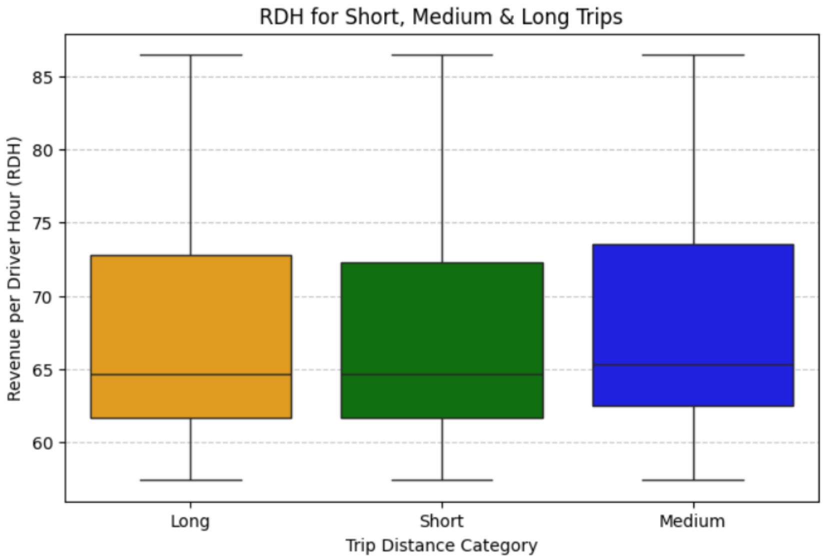

- **Finding: Trip Distance Category Is Not a Strong Predictor of RDH** The median Revenue per Driver Hour (RDH) is remarkably similar across Short, Medium, and Long trip categories. This indicates that the simple distance category of a trip is not a reliable predictor of its hourly profitability, suggesting that other factors like trip frequency and idle time are more influential.
- **Recommendation: Focus on Minimizing Idle Time** Since all trip categories offer similar earning efficiency, the algorithm's primary goal should be to minimize driver idle time by immediately queuing the *nearest* available trip for a driver, regardless of its distance category.

6. Hypothesis: Analysis of RDH by Tip Category

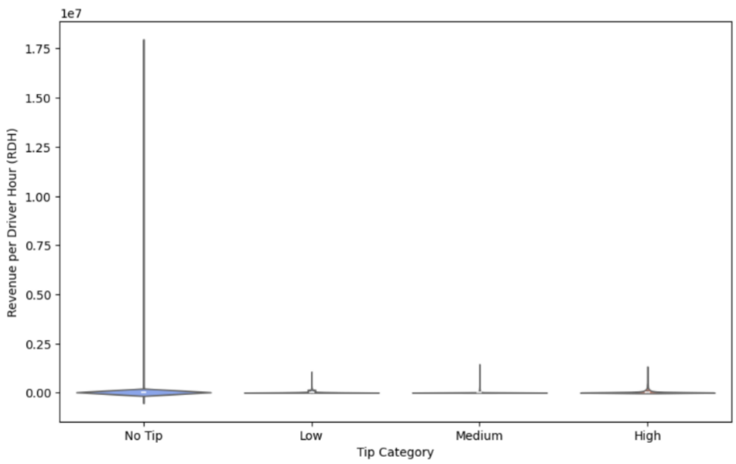

- **Finding: Data Skewed by Extreme Outliers** This analysis is heavily compromised by extreme and unrealistic outliers, particularly in the "No Tip" category where RDH values are impossibly high. These outliers completely distort the visualization, making it impossible to draw meaningful conclusions about the typical relationship between tipping and driver profitability from this chart.

- **Recommendation: Implement Aggressive Outlier Removal** Before any further analysis, a robust outlier removal process must be applied to the RDH metric. Capping the RDH at a realistic maximum (e.g., the 99th percentile) is an essential first step to create a usable and interpretable dataset.

Conclusion & Strategic Recommendations for NYC

This analysis of Uber's NYC data reveals that driver profitability (RDH) is not driven by high demand or fares, but by operational efficiency in a complex market. The key insight is that high-demand commutes often yield the lowest RDH due to gridlock, while low-supply early mornings are the most profitable. Strategy must therefore shift from managing volume to optimizing for the factors that truly impact hourly earnings, such as traffic, idle time, and supply/demand balance across the five boroughs.

Short-Term Recommendations (Immediate Actions)

- **Implement Targeted Driver Guidance:** Immediately communicate profitability insights to drivers, promoting "split-shifts" for high-RDH hours and offering "Power Zone" bonuses for specific lucrative areas like airports.
- **Conduct Actionable Geospatial Analysis:** Prioritize mapping numerical location IDs to real-world names, then analyze the key asymmetric corridors to enable predictive driver positioning along major commuter routes.

Long-Term Strategic Initiatives

- **Evolve the Matchmaking Algorithm for Profitability:** Shift the algorithm's core objective from minimizing wait time to maximizing a driver's predicted RDH, factoring in traffic, tolls, and efficient trip sequencing.
- **Restructure the Fare Model for NYC Complexity:** Strategically adjust the fare model to better compensate drivers for NYC-specific costs like tolls, congestion pricing, and unpaid return miles from outer-borough trips.

By implementing these short-term actions and investing in these long-term strategies, Uber can build a more resilient and profitable NYC operation. This data-driven approach will not only enhance driver earnings and retention in a fiercely competitive market but also solidify Uber's market leadership by creating a more efficient and reliable service for all New Yorkers.